

Lecture: Introduction to MongoDB

- What is MongoDB?
- Key Features of MongoDB
- Installation and Setup
- Basic CRUD Operations
- Data Modeling in MongoDB
- Indexing and Query Optimization
- Aggregation Framework
- Replica Sets and Sharding
- MongoDB Atlas: Managed Cloud Database
- Best Practices and Tips



Introduction to MongoDB

- MongoDB is a NoSQL database management system.
- Developed by MongoDB Inc.
- Designed for flexibility, scalability, and performance.
- Uses a document-oriented data model (documents are objects)
- Ideal for applications with rapidly changing schema and large amounts of data.



Key Features of MongoDB

- Flexible Schema: MongoDB stores data in flexible, JSON-like documents.
- Scalability: Horizontal scalability with automatic sharding.
- High Performance: Supports indexes, aggregations, and native replication.
- Rich Query Language: Supports rich queries and dynamic updates.
- High Availability: Replication with automatic failover.



- Every row in a collection does not have to be the same
- Documents (objects) may contain arrays and other documents
- Collection can easily represent a hierarchy of data
- This flexibility comes at a price in performance



Installation and Setup

- Download MongoDB Community Edition
- Install MongoDB
- Start MongoDB Server: After installation, start the MongoDB server using the appropriate command for your OS.
- Verify Installation: Open a command prompt and type 'mongo --version' to verify that MongoDB is installed correctly.
- Run mongosh to connect to the local database by default



Collection:

- In MongoDB, a collection is a group of documents.
- Analogous to a table in relational databases.
- Collections do not enforce a schema.
- Documents within a collection can have different fields and structure.



Basic CRUD Operations

- CRUD stands for Create, Read, Update, and Delete.
- These operations are fundamental for interacting with data in MongoDB.
- MongoDB provides simple and powerful methods for performing CRUD operations.
- Examples of CRUD operations include:
 - Inserting documents into a collection (Create).
 - Retrieving documents from a collection (Read).
 - Updating existing documents in a collection (Update).
 - Removing documents from a collection (Delete).



Basic CRUD Operations: Example

Create (Insert):

- To insert a document into a collection named "users":
- `db.users.insertOne({name: "John", age: 30})`

Read (Retrieve):

- To retrieve all documents from the "users" collection:
- `db.users.find()`



Data Modeling in MongoDB: Overview

- Document Structure and Schema Design
- Embedded Documents vs. References
- One-to-One Relationships
- One-to-Many Relationships
- Many-to-Many Relationships
- Denormalization and Performance
- Indexing Strategies
- Case Studies
- Best Practices



Introduction: Data Modeling

- Process of designing the structure and organization of data within MongoDB databases.
- Determine how data will be stored, accessed, and manipulated to meet the requirements of the application.
- Effective data modeling is crucial for optimizing performance, ensuring data integrity, and facilitating scalability.
- This presentation will cover various aspects of data modeling in MongoDB, including document structure, relationships, denormalization, indexing strategies, and best practices.



MongoDB Primitive Data Types

- String
- Integer
- Double
- Decimal
- BigNumber
- BigDecimal
- Boolean
- Date
- Null



MongoDB Compound Data Types

- Object
- Array
- ObjectId
- Binary Data
- Regular Expression



Examples: Scalar Data Types

```
1 {  
2   "string": "Hello, MongoDB!",  
3   "integer": 42,  
4   "double": 3.14,  
5   "boolean": true,  
6   "date": ISODate("2022-04-05T12:00:00Z"),  
7   "null_value": null  
8 }
```



Example: Integer Values

MongoDB integers are 32-bit signed

```
1 {  
2   "small_integer": 42,  
3   "large_integer": 2147483647,  
4   "negative_integer": -123  
5 }
```

Range: -2,147,483,648 to 2,147,483,647



Example: Double Values and Range

```
1 {  
2   "small_double": 3.14,  
3   "large_double": 1.7976931348623157e+308,  
4   "negative_double": -123.456  
5 }
```

Range: $\pm 5.0 \times 10^{-324}$ to $\pm 1.7 \times 10^{308}$



Example: Date Values

Dates are internally stored as a 64-bit integer

```
1 {  
2   "current_date": ISODate(),  
3   "date_with_time": ISODate("2022-04-05T12:00:00Z"),  
4   "date_only": ISODate("2022-04-05"),  
5   "date_with_timezone": ISODate("2022-04-05T00  
        :00:00-04:00")  
6 }
```



Big Integer Types

MongoDB supports two big integer types:

- **NumberLong**: A 64-bit signed integer type, capable of storing integers larger than the 32-bit signed integer range.
- **NumberDecimal**: A decimal128 type, capable of representing high-precision decimal numbers.

Usage:

- Use `NumberLong` to store integers larger than 32 bits, such as IDs, counters, or large numerical values.
- Use `NumberDecimal` for high-precision decimal calculations or when accuracy is critical.



Example: NumberLong

NumberLong is a 64-bit integer

```
1 {  
2   "user_id": NumberLong("1234567890123456789"),  
3   "large_value": NumberLong("9223372036854775807"),  
4 }
```



Document Structure and Schema Design

- Document Structure:
 - MongoDB stores data in flexible JSON-like documents called BSON.
 - Documents represent records in a collection and can vary in structure.
 - Fields within documents can be of any data type, including arrays and nested documents.
- Schema Design:
 - Designing an appropriate schema is crucial for efficient data storage and retrieval.
 - Considerations include data access patterns, query performance, and scalability.
 - Schema design decisions impact read and write operations, as well as database performance.



Example: Simple BSON Document

```
1 {  
2   "name": "John Doe",  
3   "age": 30,  
4   "email": "john@example.com"  
5 }
```



Example: Nested BSON Document

```
1 {  
2   "name": "John Doe",  
3   "age": 30,  
4   "address": {  
5     "city": "New York",  
6     "zip": "10001"  
7   }  
8 }
```



Example: Array in BSON Document

```
1 {  
2   "name": "John Doe",  
3   "age": 30,  
4   "hobbies": ["reading", "hiking", "photography"]  
5 }
```



Object Type

- Represents a nested document within a MongoDB document.
- Allows for the grouping of related fields and values.
- Can be used to model complex data structures.
- Supports nested objects within objects (subdocuments).



Example: Object Type

```
1 {  
2   "student": {  
3     "name": "Alice",  
4     "age": 25,  
5     "grades": [90, 85, 95],  
6     "address": {  
7       "city": "New York",  
8       "zipcode": "10001"  
9     }  
10  }  
11 }
```



Collection Naming Conventions

- Use plural nouns
- Be descriptive
- Use camelCase or underscores
- Avoid reserved characters \$, ., 0
- Keep it short and concise
- Maintain consistency



Collection Naming: Legal and Illegal Examples

- users
- blog_posts
- customerFeedback
- userAccounts
- \$users
- my.collection
- user.name
- text
0me
- collection1
- 1collection



Insert Examples

```
1 // Create the "users" collection
2 db.createCollection("users")
3
4 // Insert a single document using insertOne()
5 db.users.insertOne({
6     username: "john_doe",
7     email: "john@example.com",
8     age: 30
9 })
```



Example: insertMany

```
1 // Insert multiple documents using insertMany()
2 db.users.insertMany([
3 {
4     username: "jane_doe",
5     email: "jane@example.com",
6     age: 25
7 },
8 {
9     username: "bob_smith",
10    email: "bob@example.com",
11    age: 35
12 }
13 ])
```



Create Collection with Validation

```
1 // Create the foodSubstances collection with validation
2 db.createCollection("foodSubstances", {
3   validator: {
4     $jsonSchema: {
5       bsonType: "object",
6       required: ["name", "type", "altname", "unit"],
7       properties: {
8         name: { bsonType: "string" },
9         type: { enum: ["vitamin", "mineral", "poison", "
10           allergen", "drug"] },
11         altname: { bsonType: ["string", "null"] },
12         unit: { enum: ["mg", "mcg", "IU", "g", "ug"] }
13       }
14     }
15 });
```



Inserting Records with Validation

```
1 // Inserting a correct record (Vitamin C)
2 db.foodSubstances.insertOne
3 ({name: "vC", type: "vitamin", altname: "Ascorbic Acid",
   unit: "mg"});
4
5 // Inserting an incorrect record (violates constraint)
6 db.foodSubstances.insertOne
7 ({name: "mFe", type: "foo", altname: "Iron", unit: "mg"
   });
8
9 db.foodSubstances.insertOne
0 ({name: "vC", type: "vitamin", altname: "Ascorbic Acid",
   unit: "mg"});
1 // error, vitamin C already exists (no constraint
   prevents this)
```



MongoDB Comparison Operators

- \$eq - Matches values that are equal to a specified value.
- \$ne - Matches values that are not equal to a specified value.
- \$gt - Matches values that are greater than a specified value.
- \$gte - Matches values that are greater than or equal to a specified value.
- \$lt - Matches values that are less than a specified value.
- \$lte - Matches values that are less than or equal to a specified value.
- \$in - Matches any of the values specified in an array.
- \$nin - Matches none of the values specified in an array.



Examples of Comparison Operations

```
1 // Find documents where the "age" field is equal to 30
2 db.collection.find({ age: { $eq: 30 } })
3
4 // Find documents where the "price" field is greater
   than $100
5 db.collection.find({ price: { $gt: 100 } })
6
7 // Find documents where the "quantity" field is less
   than or equal to 10
8 db.collection.find({ quantity: { $lte: 10 } })
```



Examples of Array Operations

```
1 // Add a grade of 85 to Melanie's grades array
2 db.students.updateOne( { _id: 1 }, { $push: { grades: 85
  } })
3
4 // Add a grade of 90 if it doesn't already exist, note
  creates hashset
5 db.students.updateOne( { _id: 1 }, { $addToSet: { grades:
  90 } })
6
7 // Remove the last grade from Melanie's grades array
8 db.students.updateOne( { _id: 1 }, { $pop: { grades: 1 }
  } )
9
0 // Remove all grades less than 80 from Melanie's grades
  array
1 db.students.updateOne( { _id: 1 }, { $pull: { grades: {
  $lt: 80 } } } )
2
3 // Remove grades 85 and 90 from Melanie's grades array
4 db.students.updateOne( { _id: 1 }, { $pullAll: { grades:
  [85, 90] } } )
```

Examples of Regex Queries in MongoDB

```
1 // Finding a student whose name starts with "E"
2 db.students.find({ name: /^E/ })
3
4 // Finding students whose name contains "a", "b", or "c"
5 db.students.find({ name: /[abc]/ })
6
7 // Finding students whose name ends with "y"
8 db.students.find({ name: /y$/ })
9
10 // Finding students whose name contains "mel" case
    insensitively
11 db.students.find({ name: /mel/i })
```



Storing JavaScript in a Database

- **Custom Business Logic:** Store validation rules, workflow logic, or calculation formulas as JavaScript functions in the database.
- **Dynamic Scripting:** Dynamically generate and execute scripts based on user input or configuration for flexible application behavior.
- **Stored Procedures:** Define complex database operations as JavaScript functions for server-side execution.
- **Content Management Systems (CMS):** Store templates, themes, or dynamic content logic as JavaScript code for easy customization.
- **User-defined Functions:** Allow users to define custom functions or scripts, stored in the database, for extended application functionality.



Example Business Logic

```
1 // Example business logic: Calculate total price with  
  discount  
2 function calculateTotalPrice(price, discount) {  
3     return price - (price * discount);  
4 }
```



- BTree is a popular indexing type for databases
- BTrees allow sorting data
- Hashing is faster, but only allows equal/not equal
- Hashing can only be used where ordering is not relevant
- Hashing floating point is weird because of roundoff error

BTree



Indexing:

- MongoDB supports various types of indexes to optimize query performance.
- Common index types include single-field, compound, multi-key, text, and geospatial indexes.
- Indexes help MongoDB efficiently locate and retrieve documents based on query conditions.

Query Optimization:

- Effective query optimization is essential for improving database performance.
- Techniques include utilizing indexes, minimizing the number of documents scanned, and optimizing query execution plans.
- Understanding query patterns and access patterns is crucial for identifying optimization opportunities.



- Ascending/Descending Order:
 - Specifies the sorting order of the index.
 - 1 indicates ascending order, -1 indicates descending order.
- Unique:
 - Ensures that indexed field(s) contain unique values.
 - MongoDB rejects operations that would result in duplicate values.
- Not Null Constraint:
 - Enforced by the validator, not the index.
 - Validators specify rules for document validation, including required fields.



Creating Unique Index on Short Name

```
1 // Create unique index on the shortName field
2 db.foodSubstances.createIndex({ name: 1 }, { unique:
    true });
```



Example: Index Substances by Type

To efficiently query substances by their type, create an index on the type field in the foodSubstances collection:

```
1 db.foodSubstances.createIndex({ "type": 1 });  
2 db.foodSubstances.createIndex({ "type": "hashed" })
```

Example query to retrieve all substances of a specific type:

```
1 db.foodSubstances.find({ "type": "vitamin" });
```



Updating Food Substances Collection

```
1 db.foodSubstances.updateOne(  
2   { name: "vC" },  
3   { $set: { "RDA.child<12": "25-75" } }  
4 );  
5 db.foodSubstances.updateMany(  
6   { type: "vitamin" }, // Filter: Update all documents  
7   { $set: { "RDA.child<12": "30-70" } } // syntax  
8   probably wrong. I should not have used < in  
9   string  
10 );
```



- High availability and fault tolerance
- Primary and secondary nodes
- Automatic failover
- Read scalability
- Ensures data durability and minimizes downtime



Updating Food Substances Collection

```
1 # Start MongoDB instances on different servers
2 mongod --port 27017 --dbpath /data/db1 --replSet rs0
3 mongod --port 27018 --dbpath /data/db2 --replSet rs0
4 mongod --port 27019 --dbpath /data/db3 --replSet rs0
5
6 # Connect to one of the MongoDB instances
7 mongo --port 27017
8
9 # Initialize the replica set configuration
10 rs.initiate({_id: "rs0", members: [{_id: 0, host: "
    localhost:27017"}]})
11
12 # Add the other MongoDB instances to the replica set
13 rs.add("localhost:27018")
14 rs.add("localhost:27019")
```



Sharding Overview

- Horizontal scaling technique
- Distributes data across shards
- A shard is a subset of data
- Each shard can run on a separate MongoDB instance
- Net performance increases by parallelizing MongoDB
- Data distribution based on shard key
- Enables scalability



Horizontal Sharding

- Splitting data horizontally involves dividing the entire dataset into smaller subsets or partitions.
- Each partition is stored on a separate server.
- This approach is useful for distributing large datasets across multiple servers to improve performance and scalability.



Horizontal Sharding Example

- Sharded Collection: "users"
- Shard Key: "country"
- Servers:
 - S1: Stores users from USA, Canada, Mexico
 - S2: Stores users from UK, Germany, France
 - S3: Stores users from Japan, China, India



- Splitting data vertically involves dividing the columns of a table across different servers.
- Each server holds a subset of columns for all rows in the table.
- This approach is useful for separating frequently accessed columns or those requiring more resources.



- MongoDB is a powerful and flexible
- Distributed database, easy to use
- Slower than mysql, easy parallelism might make up for that

